

Lab 8

#Breast Cancer Project Today we are going to explore some data from the Univeristy of Wisconsin Cancer Center on Breast biosy data.

```
wisc.data <-read.csv("/Users/JoshTran/Desktop/BIMM 143/Week 5/WisconsinCancer.csv", row.n
head(wisc.data)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1
84358402	M	20.29	14.34	135.10	1297.0
843786	M	12.45	15.70	82.57	477.1
	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean	
842302	0.11840	0.27760	0.3001	0.14710	
842517	0.08474	0.07864	0.0869	0.07017	
84300903	0.10960	0.15990	0.1974	0.12790	
84348301	0.14250	0.28390	0.2414	0.10520	
84358402	0.10030	0.13280	0.1980	0.10430	
843786	0.12780	0.17000	0.1578	0.08089	
	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.2419	0.07871	1.0950	0.9053	8.589
842517	0.1812	0.05667	0.5435	0.7339	3.398
84300903	0.2069	0.05999	0.7456	0.7869	4.585
84348301	0.2597	0.09744	0.4956	1.1560	3.445
84358402	0.1809	0.05883	0.7572	0.7813	5.438
843786	0.2087	0.07613	0.3345	0.8902	2.217
	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se
842302	153.40	0.006399	0.04904	0.05373	0.01587
842517	74.08	0.005225	0.01308	0.01860	0.01340
84300903	94.03	0.006150	0.04006	0.03832	0.02058
84348301	27.23	0.009110	0.07458	0.05661	0.01867
84358402	94.44	0.011490	0.02461	0.05688	0.01885
843786	27.19	0.007510	0.03345	0.03672	0.01137
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst	
842302	0.03003	0.006193	25.38	17.33	
842517	0.01389	0.003532	24.99	23.41	
84300903	0.02250	0.004571	23.57	25.53	
84348301	0.05963	0.009208	14.91	26.50	
84358402	0.01756	0.005115	22.54	16.67	
843786	0.02165	0.005082	15.47	23.75	
	perimeter_worst	area_worst	smoothness_worst	compactness_worst	
842302	184.60	2019.0	0.1622	0.6656	
842517	158.80	1956.0	0.1238	0.1866	
84300903	152.50	1709.0	0.1444	0.4245	
84348301	98.87	567.7	0.2098	0.8663	
84358402	152.20	1575.0	0.1374	0.2050	

843786	103.40	741.6	0.1791	0.5249
	concavity_worst	concave.points_worst	symmetry_worst	
842302	0.7119	0.2654	0.4601	
842517	0.2416	0.1860	0.2750	
84300903	0.4504	0.2430	0.3613	
84348301	0.6869	0.2575	0.6638	
84358402	0.4000	0.1625	0.2364	
843786	0.5355	0.1741	0.3985	
	fractal_dimension_worst			
842302	0.11890			
842517	0.08902			
84300903	0.08758			
84348301	0.17300			
84358402	0.07678			
843786	0.12440			

Q1. How many patients samples are in this dataset?

```
nrow(wisc.data)
```

```
[1] 569
```

There are 569 patients in this data set.

Q2. How many cancer (M) and non cancer (B) samples are there?

```
table(wisc.data$diagnosis)
```

```

B    M
357 212

```

```
head(wisc.data)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1
84358402	M	20.29	14.34	135.10	1297.0
843786	M	12.45	15.70	82.57	477.1
	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean	
842302	0.11840	0.27760	0.3001	0.14710	
842517	0.08474	0.07864	0.0869	0.07017	
84300903	0.10960	0.15990	0.1974	0.12790	
84348301	0.14250	0.28390	0.2414	0.10520	
84358402	0.10030	0.13280	0.1980	0.10430	
843786	0.12780	0.17000	0.1578	0.08089	

	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.2419	0.07871	1.0950	0.9053	8.589
842517	0.1812	0.05667	0.5435	0.7339	3.398
84300903	0.2069	0.05999	0.7456	0.7869	4.585
84348301	0.2597	0.09744	0.4956	1.1560	3.445
84358402	0.1809	0.05883	0.7572	0.7813	5.438
843786	0.2087	0.07613	0.3345	0.8902	2.217

	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se
842302	153.40	0.006399	0.04904	0.05373	0.01587
842517	74.08	0.005225	0.01308	0.01860	0.01340
84300903	94.03	0.006150	0.04006	0.03832	0.02058
84348301	27.23	0.009110	0.07458	0.05661	0.01867
84358402	94.44	0.011490	0.02461	0.05688	0.01885
843786	27.19	0.007510	0.03345	0.03672	0.01137

	symmetry_se	fractal_dimension_se	radius_worst	texture_worst
842302	0.03003	0.006193	25.38	17.33
842517	0.01389	0.003532	24.99	23.41
84300903	0.02250	0.004571	23.57	25.53
84348301	0.05963	0.009208	14.91	26.50
84358402	0.01756	0.005115	22.54	16.67
843786	0.02165	0.005082	15.47	23.75

	perimeter_worst	area_worst	smoothness_worst	compactness_worst
842302	184.60	2019.0	0.1622	0.6656
842517	158.80	1956.0	0.1238	0.1866
84300903	152.50	1709.0	0.1444	0.4245
84348301	98.87	567.7	0.2098	0.8663
84358402	152.20	1575.0	0.1374	0.2050
843786	103.40	741.6	0.1791	0.5249

	concavity_worst	concave.points_worst	symmetry_worst
842302	0.7119	0.2654	0.4601
842517	0.2416	0.1860	0.2750
84300903	0.4504	0.2430	0.3613
84348301	0.6869	0.2575	0.6638
84358402	0.4000	0.1625	0.2364
843786	0.5355	0.1741	0.3985

	fractal_dimension_worst
842302	0.11890
842517	0.08902
84300903	0.08758
84348301	0.17300
84358402	0.07678
843786	0.12440

```
diagnosis <- as.factor(wisc.data$diagnosis)
#diagnosis
```

```
head(diagnosis)
```

```
[1] M M M M M M
```

Levels: B M

Now exclude the diagnosis cloumn from the data

```
wisc.data <- wisc.data[,-1]
```

Q3. How many "dimensions","Variables", "columns", are in this dataset?

```
ncol(wisc.data)
```

```
[1] 30
```

#Principal Component Analysis (PCA)

To perform PCA in R we can use `prcomp()` function. it takes as input a numeric dataset and optional `scale = FALSE/TRUE` argument.

We generally always want to set the `scale=TRUE` but lets make sure by checking if the mean and standard deviation values are different across these 30 columns.

```
round(colMeans(wisc.data) )
```

radius_mean	texture_mean	perimeter_mean
14	19	92
area_mean	smoothness_mean	compactness_mean
655	0	0
concavity_mean	concave.points_mean	symmetry_mean
0	0	0
fractal_dimension_mean	radius_se	texture_se
0	0	1
perimeter_se	area_se	smoothness_se
3	40	0
compactness_se	concavity_se	concave.points_se
0	0	0
symmetry_se	fractal_dimension_se	radius_worst
0	0	16
texture_worst	perimeter_worst	area_worst
26	107	881
smoothness_worst	compactness_worst	concavity_worst
0	0	0
concave.points_worst	symmetry_worst	fractal_dimension_worst
0	0	0

```
pca <- prcomp(wisc.data, scale = TRUE)
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)? 0.4427

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data? 0.72636

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data? PC7

```
attributes(pca$x)
```

```
$dim
```

```
[1] 569 30
```

```
$dimnames
```

```
$dimnames[[1]]
```

```

[1] "842302" "842517" "84300903" "84348301" "84358402" "843786"
[7] "844359" "84458202" "844981" "84501001" "845636" "84610002"
[13] "846226" "846381" "84667401" "84799002" "848406" "84862001"
[19] "849014" "8510426" "8510653" "8510824" "8511133" "851509"
[25] "852552" "852631" "852763" "852781" "852973" "853201"
[31] "853401" "853612" "85382601" "854002" "854039" "854253"
[37] "854268" "854941" "855133" "855138" "855167" "855563"
[43] "855625" "856106" "85638502" "857010" "85713702" "85715"
[49] "857155" "857156" "857343" "857373" "857374" "857392"
[55] "857438" "85759902" "857637" "857793" "857810" "858477"
```

[61]	"858970"	"858981"	"858986"	"859196"	"85922302"	"859283"
[67]	"859464"	"859465"	"859471"	"859487"	"859575"	"859711"
[73]	"859717"	"859983"	"8610175"	"8610404"	"8610629"	"8610637"
[79]	"8610862"	"8610908"	"861103"	"8611161"	"8611555"	"8611792"
[85]	"8612080"	"8612399"	"86135501"	"86135502"	"861597"	"861598"
[91]	"861648"	"861799"	"861853"	"862009"	"862028"	"86208"
[97]	"86211"	"862261"	"862485"	"862548"	"862717"	"862722"
[103]	"862965"	"862980"	"862989"	"863030"	"863031"	"863270"
[109]	"86355"	"864018"	"864033"	"86408"	"86409"	"864292"
[115]	"864496"	"864685"	"864726"	"864729"	"864877"	"865128"
[121]	"865137"	"86517"	"865423"	"865432"	"865468"	"86561"
[127]	"866083"	"866203"	"866458"	"866674"	"866714"	"8670"
[133]	"86730502"	"867387"	"867739"	"868202"	"868223"	"868682"
[139]	"868826"	"868871"	"868999"	"869104"	"869218"	"869224"
[145]	"869254"	"869476"	"869691"	"86973701"	"86973702"	"869931"
[151]	"871001501"	"871001502"	"8710441"	"87106"	"8711002"	"8711003"
[157]	"8711202"	"8711216"	"871122"	"871149"	"8711561"	"8711803"
[163]	"871201"	"8712064"	"8712289"	"8712291"	"87127"	"8712729"
[169]	"8712766"	"8712853"	"87139402"	"87163"	"87164"	"871641"
[175]	"871642"	"872113"	"872608"	"87281702"	"873357"	"873586"
[181]	"873592"	"873593"	"873701"	"873843"	"873885"	"874158"
[187]	"874217"	"874373"	"874662"	"874839"	"874858"	"875093"
[193]	"875099"	"875263"	"87556202"	"875878"	"875938"	"877159"
[199]	"877486"	"877500"	"877501"	"877989"	"878796"	"87880"
[205]	"87930"	"879523"	"879804"	"879830"	"8810158"	"8810436"
[211]	"881046502"	"8810528"	"8810703"	"881094802"	"8810955"	"8810987"
[217]	"8811523"	"8811779"	"8811842"	"88119002"	"8812816"	"8812818"
[223]	"8812844"	"8812877"	"8813129"	"88143502"	"88147101"	"88147102"
[229]	"88147202"	"881861"	"881972"	"88199202"	"88203002"	"88206102"
[235]	"882488"	"88249602"	"88299702"	"883263"	"883270"	"88330202"
[241]	"88350402"	"883539"	"883852"	"88411702"	"884180"	"884437"
[247]	"884448"	"884626"	"88466802"	"884689"	"884948"	"88518501"
[253]	"885429"	"8860702"	"886226"	"886452"	"88649001"	"886776"
[259]	"887181"	"88725602"	"887549"	"888264"	"888570"	"889403"
[265]	"889719"	"88995002"	"8910251"	"8910499"	"8910506"	"8910720"
[271]	"8910721"	"8910748"	"8910988"	"8910996"	"8911163"	"8911164"
[277]	"8911230"	"8911670"	"8911800"	"8911834"	"8912049"	"8912055"
[283]	"89122"	"8912280"	"8912284"	"8912521"	"8912909"	"8913"
[289]	"8913049"	"89143601"	"89143602"	"8915"	"891670"	"891703"
[295]	"891716"	"891923"	"891936"	"892189"	"892214"	"892399"
[301]	"892438"	"892604"	"89263202"	"892657"	"89296"	"893061"
[307]	"89344"	"89346"	"893526"	"893548"	"893783"	"89382601"
[313]	"89382602"	"893988"	"894047"	"894089"	"894090"	"894326"
[319]	"894329"	"894335"	"894604"	"894618"	"894855"	"895100"
[325]	"89511501"	"89511502"	"89524"	"895299"	"8953902"	"895633"
[331]	"896839"	"896864"	"897132"	"897137"	"897374"	"89742801"
[337]	"897604"	"897630"	"897880"	"89812"	"89813"	"898143"
[343]	"89827"	"898431"	"89864002"	"898677"	"898678"	"89869"
[349]	"898690"	"899147"	"899187"	"899667"	"899987"	"9010018"
[355]	"901011"	"9010258"	"9010259"	"901028"	"9010333"	"901034301"
[361]	"901034302"	"901041"	"9010598"	"9010872"	"9010877"	"901088"

```
[367] "9011494" "9011495" "9011971" "9012000" "9012315" "9012568"
[373] "9012795" "901288" "9013005" "901303" "901315" "9013579"
[379] "9013594" "9013838" "901549" "901836" "90250" "90251"
[385] "902727" "90291" "902975" "902976" "903011" "90312"
[391] "90317302" "903483" "903507" "903516" "903554" "903811"
[397] "90401601" "90401602" "904302" "904357" "90439701" "904647"
[403] "904689" "9047" "904969" "904971" "905189" "905190"
[409] "90524101" "905501" "905502" "905520" "905539" "905557"
[415] "905680" "905686" "905978" "90602302" "906024" "906290"
[421] "906539" "906564" "906616" "906878" "907145" "907367"
[427] "907409" "90745" "90769601" "90769602" "907914" "907915"
[433] "908194" "908445" "908469" "908489" "908916" "909220"
[439] "909231" "909410" "909411" "909445" "90944601" "909777"
[445] "9110127" "9110720" "9110732" "9110944" "911150" "911157302"
[451] "9111596" "9111805" "9111843" "911201" "911202" "9112085"
[457] "9112366" "9112367" "9112594" "9112712" "911296201" "911296202"
[463] "9113156" "911320501" "911320502" "9113239" "9113455" "9113514"
[469] "9113538" "911366" "9113778" "9113816" "911384" "9113846"
[475] "911391" "911408" "911654" "911673" "911685" "911916"
[481] "912193" "91227" "912519" "912558" "912600" "913063"
[487] "913102" "913505" "913512" "913535" "91376701" "91376702"
[493] "914062" "914101" "914102" "914333" "914366" "914580"
[499] "914769" "91485" "914862" "91504" "91505" "915143"
[505] "915186" "915276" "91544001" "91544002" "915452" "915460"
[511] "91550" "915664" "915691" "915940" "91594602" "916221"
[517] "916799" "916838" "917062" "917080" "917092" "91762702"
[523] "91789" "917896" "917897" "91805" "91813701" "91813702"
[529] "918192" "918465" "91858" "91903901" "91903902" "91930402"
[535] "919537" "919555" "91979701" "919812" "921092" "921362"
[541] "921385" "921386" "921644" "922296" "922297" "922576"
[547] "922577" "922840" "923169" "923465" "923748" "923780"
[553] "924084" "924342" "924632" "924934" "924964" "925236"
[559] "925277" "925291" "925292" "925311" "925622" "926125"
[565] "926424" "926682" "926954" "927241" "92751"
```


```
$dimnames[[2]]
```

```
[1] "PC1" "PC2" "PC3" "PC4" "PC5" "PC6" "PC7" "PC8" "PC9" "PC10"
[11] "PC11" "PC12" "PC13" "PC14" "PC15" "PC16" "PC17" "PC18" "PC19" "PC20"
[21] "PC21" "PC22" "PC23" "PC24" "PC25" "PC26" "PC27" "PC28" "PC29" "PC30"
```

Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

This is a hot mess of a plot and we will need to generate our own plots to make sense of this PCA result.


```
plot(pca$x[,1], pca$x[,2], col= diagnosis)
```



```
# Create a data.frame for ggplot
df <- as.data.frame(pca$x)
df$diagnosis <- diagnosis

# Load the ggplot2 package
library(ggplot2)

# Make a scatter plot colored by diagnosis
ggplot(df) +
  aes(PC1, PC2, col= diagnosis) +
  geom_point()
```

Q8. How much variance is captured in the top 3 PCs.

They capture 72.636% of the total variance.

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`?

This tells us how much this original feature contributes to the first PC.

```
pca$rotation["concave.points_mean",1]
```

```
[1] -0.2608538
```

```
attributes(pca)
```

```
$names
```

```
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

```
$class
```

```
[1] "prcomp"
```

#Combine PCA results with clustering

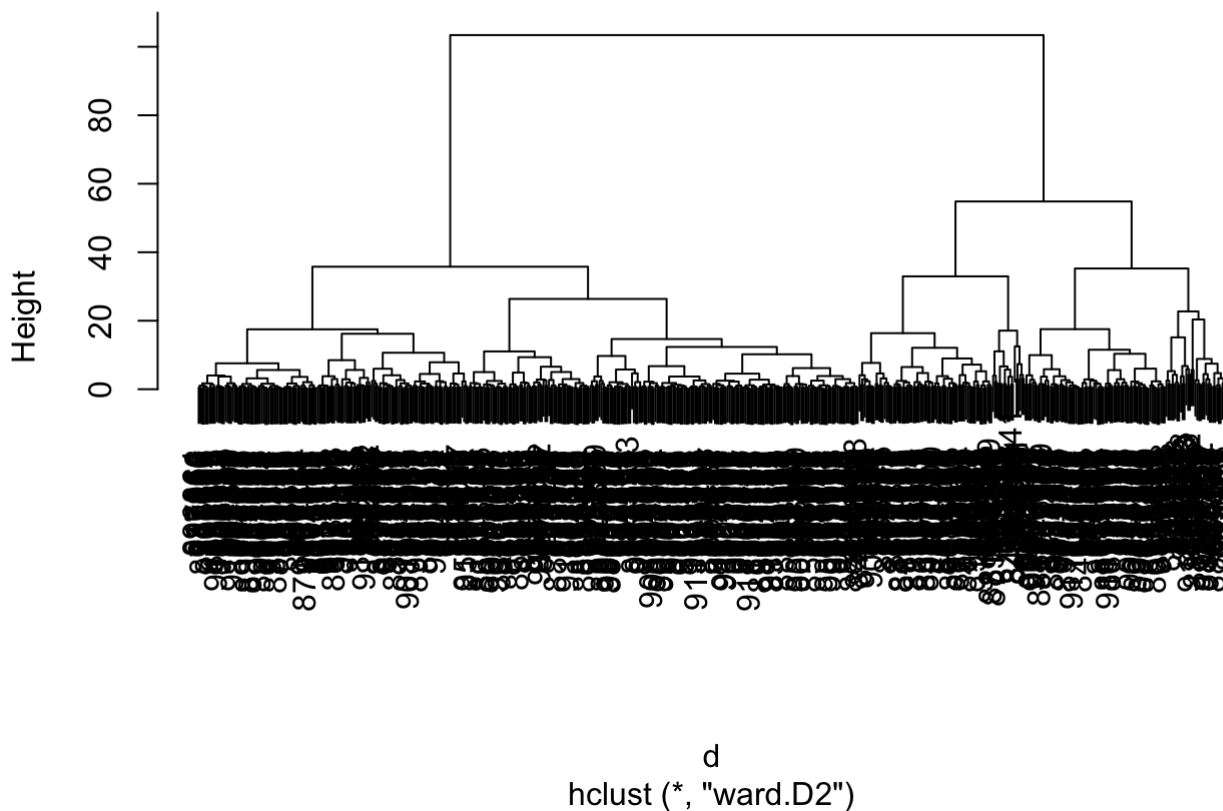
We can use our new PCA variables (i.e the scores along the PCs contained in R `pca$x`) as input for other methods such as clustering

#Hclust needs a distance matrix as input

```
d <- dist(pca$x[,1:3])
```

```
hc <- hclust(d, method="ward.D2")
plot(hc)
```

Cluster Dendrogram



Q11. Using the `plot()` and `abline()` functions, what is the height at which the clustering model has 4 clusters?

`h = 35.5`

To get our cluster membership vector we can use the `cutree()` function and specify a height (`h`) or number of groups (`k`).

```
grps <- cutree(hc, h=35.5)
table(grps)
```

```
grps
 1  2  3  4
```

```
111  92 216 150
```

I want to find out how many diagnosis "M" and "B" are in each grp?

```
table(diagnosis)
```

```
diagnosis
```

```
  B   M  
357 212
```

Q12. Can you find a better cluster vs diagnoses match by cutting into a different number of clusters between 2 and 10?


```
table(diagnosis, grps)
```

```
      grps  
diagnosis  1   2   3   4  
  B    0  24 184 149  
  M 111  68  32   1
```

Q13. Which method gives your favorite results for the same data.dist dataset? Explain your reasoning.

We can also plot our results using our clustering vector `grps`


```
plot(pca$x[,1:2], col=grps)
```



```
# Create a data.frame for ggplot
df <- as.data.frame(pca$x)
df$diagnosis <- diagnosis

# Load the ggplot2 package
library(ggplot2)

# Make a scatter plot colored by diagnosis
ggplot(df) +
  aes(PC1, PC2, col= grps) +
  geom_point()
```



Q. Q15. What is the specificity ($TN/(TN+FN)$) and sensitivity($TP/(TP+FN)$) of our current results?

Specificity= $(333/(333+33)) = .91$ Sensitivity = $(179/(179+33)) = .84$

Prediction

Q16. Which of these new patients should we prioritize for follow up based on your results?

Patient 2 should follow up.